

Overview of Deep Learning

Tianxiang (Adam) Gao

September 9, 2024

Outline

- 1 What is Deep Learning?
- 2 Brief History of Neural Networks
- 3 Perceptron
- 4 Multilayer Perceptrons

Recap: What is AI?

- **Artificial Intelligence (AI)** is a broad field that focuses on creating systems capable of performing tasks that typically require human intelligence.
- To pass the *(total) Turing Test*, it needs:
 - *Natural language processing (NLP)* to enable it to communicate effectively;
 - *Knowledge representation* to store and retrieve information;
 - *Automated reasoning* to use the stored information to answer questions and to draw new conclusions;
 - *Machine learning* to recognize patterns from data and adapt to new situations;
 - *Computer vision (CV)* to perceive objects;
 - *Robotics* to manipulate and interact with the physical world.

Recap: What is ML?

- **Machine learning (ML)** is a subset of AI that focuses on developing algorithms and (statistical) models to learn from (training) data and generalize to unseen data.
- Commonly used algorithms and models are categorized into:
 - *Supervised learning*: linear/logistic regression, support vector machines (SVM), decision trees, and neural networks
 - *Unsupervised learning*: k-means clustering, dimensionality reduction, Gaussian mixture models, generative models
 - *Reinforcement learning*: Q-Learning, policy gradient, deep Q-networks (DQN)

What is DL?

- **Deep learning (DL)** is a subset of ML that focuses on using *deep neural networks (DNN)* with many layers to learn representation in large datasets.
 - It is capable of **automatically** learning features from raw data, unlike other machine learning models that rely on manually crafted features.
 - It learns the intricate structures from data by using *backpropagation* to update the parameters in DNN.
 - It encompasses various neural network architectures, including *convolutional neural Networks (CNNs)*, *recurrent neural networks (RNNs)*, *transformers*, and *graph neural networks (GNNs)*, each tailored to specific tasks.
 - It has led to significant breakthroughs in various applications, such as CV, NLP, speech recognition, and biomedical science.

- 1 What is Deep Learning?
- 2 **Brief History of Neural Networks**
- 3 Perceptron
- 4 Multilayer Perceptrons

Early Beginnings and “AI Winters”

1940s: Early Beginnings

- 1 In 1943, Warren McCulloch and Walter Pitts introduced the first mathematical model of a neuron, the **McCulloch-Pitts Neuron Model**.
- 2 In 1949, Donald Hebb proposed **Hebbian learning** in his book *The Organization of Behavior*, summarized as “cells that fire together wire together.”

1950s-1960s: Perceptron and the First “AI Winter”

- 1 In 1958, Frank Rosenblatt proposed the **perceptron**, an early neural network model.
- 2 In 1969, Marvin Minsky and Seymour Papert demonstrated that the perceptron could not solve **non-linear** problems, such as the **XOR problem**, leading to the first “AI winter.”

1980s-1990s: Revival with Expert Systems and Backpropagation

- 1 In 1980, **XCON** became one of the first commercially successful expert systems, marking a turning point for AI in the industry.
- 2 In 1986, Geoffrey Hinton, David Rumelhart, and Ronald Williams popularized **backpropagation** in their *Nature* paper, “Learning Representations by Back-Propagating Errors.”

1990s-2000s: The Second “AI Winter”

- The second “AI winter” occurred mainly due to the failure of **expert systems** to scale and generalize beyond narrow, rule-based tasks.
- Neural networks, despite the promise shown after the introduction of backpropagation, were still constrained by limited **computational power**, **scalability issues**, and **insufficient data**.

Emergence of Deep Learning

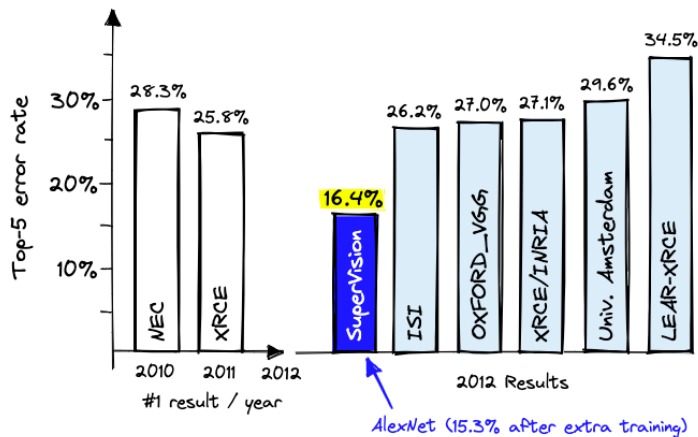
1990s-2000s: Advancements and the Emergence of Deep Learning

- In 1989, Yann LeCun et al. created the LetNet as an early example of a **CNN**, which became critical for image processing tasks.
- In 1997, the **Long Short-Term Memory (LSTM) network** was proposed by Sepp Hochreiter and Jürgen Schmidhuber that overcome the problem of vanishing gradients in RNNs

Modern Era: AlexNet 2012

2010s-Present: Deep Learning Revolution

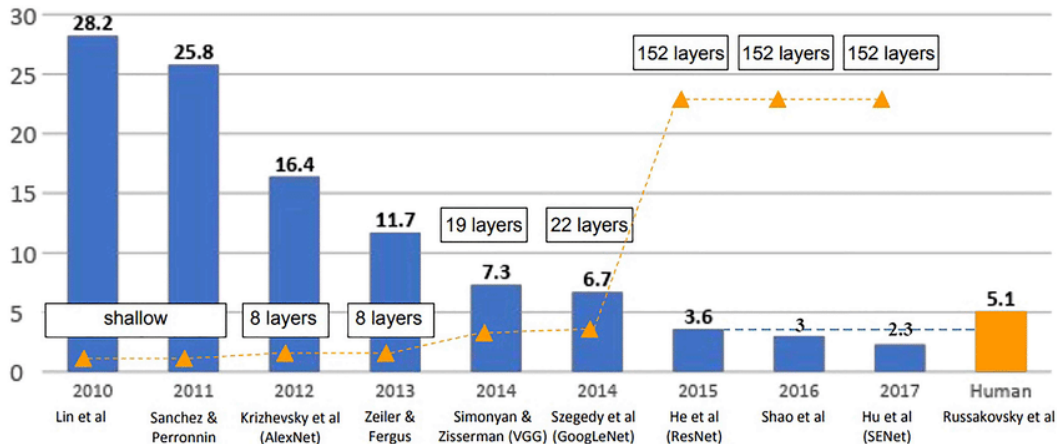
- In 2012, AlexNet trained using **GPUs** dominated the ImageNet competition



Modern Era: ResNet 2015

2010s-Present: Deep Learning Revolution

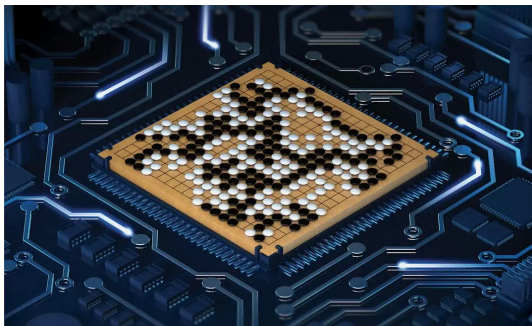
- In 2015, the introduction of skip connection in **ResNet** allowed the training of extremely ResNet, e.g., 152 layers, achieving **better-than-human** performance on ImageNet



Modern Era: AlphaGo 2016-2017

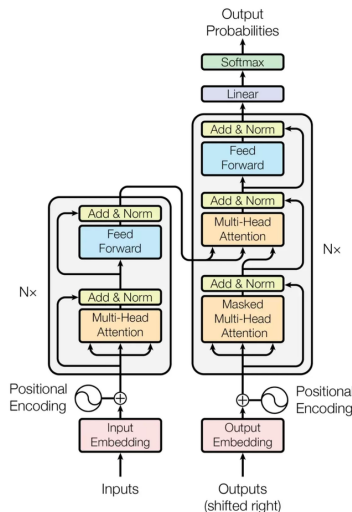
2010s-Present: Deep Learning Revolution

- In 2016, **AlphaGo** gained worldwide attention by defeating South Korean professional Go player Lee Sedol, one of the best players in the world, showing that AI can master complex and strategic games.
- In 2017, AlphaGo further cemented its dominance by defeating the world's number one Go player, Ke Jie of China, in a best-of-three match, **winning all three games**.



Modern Era: The Transformer and Attention 2017-2022

2010s-Present: Deep Learning Revolution

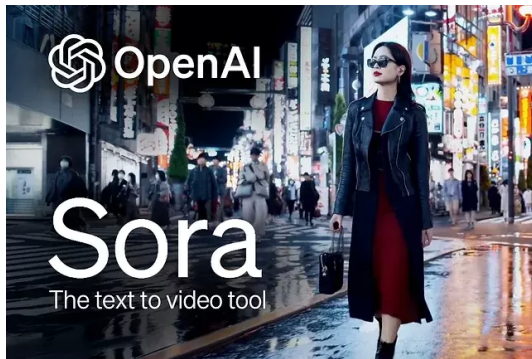


- In 2017, the **Transformer architecture** and **attention mechanism** were introduced, revolutionizing NLP with superior performance in sequence-based tasks.
- In 2018, **BERT** and **GPT** emerged as foundational pre-trained models in NLP, significantly improving performance across a wide range of downstream tasks.
- In 2020, OpenAI developed **GPT-3** as the largest language model at time with **175 billion parameters**, trained on **499 billion tokens** (approximately **570 GB of text**) using around **10,000 GPUs** over several months.
- In 2022, **ChatGPT** was released by OpenAI, based on the GPT-3.5 model. It quickly gained widespread attention for its ability to generate human-like text.

Modern Era: Generative AI 2021-Present

2010s-Present: Deep Learning Revolution

- In 2021, **Diffusion models** like Denoising Diffusion Probabilistic Models (DDPM) became prominent for generating high-quality images, challenging the dominance of GANs.
- In 2022, Stability AI released **Stable Diffusion**, allowing users to generate images from text prompts based on diffusion models.
- In 2024, OpenAI introduced **Sora**, a text-to-video model based on diffusion models that can generate about 1 minute of high-quality video from text prompts.

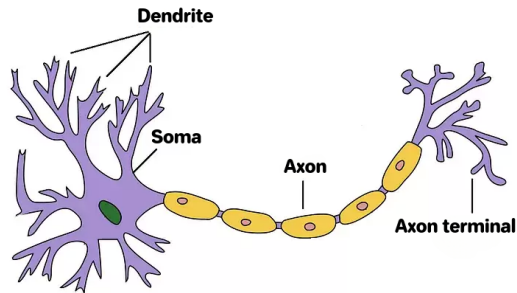


Applications of Deep Learning

- **Computer Vision**
- **Natural Language Processing (NLP)**
- **Speech Recognition and Generation**
- **Biomedical Science and Healthcare**
- Self-Driving Vehicles
- Recommendation Systems
- Finance and Fraud Detection
- Medical Imaging and Diagnostics
- Robotics and Automation
- Music Technology and Audio Processing
- Climate Science and Environmental Monitoring
- Biological Sciences and Bioinformatics
- Agricultural Technology
- Meteorology and Weather Prediction
- Cloud Computing and Data Centers
- Smart Manufacturing and Industry 4.0
- Logistics and Supply Chain Optimization
- Food Security and Sustainable Agriculture
- Cybersecurity and Threat Detection
- Software Engineering and DevOps
- Materials Science and Engineering
- Mechanical Engineering and System Design
- Digital Media, Filmmaking, and Animation
- User Interface and User Experience (UI/UX) Development

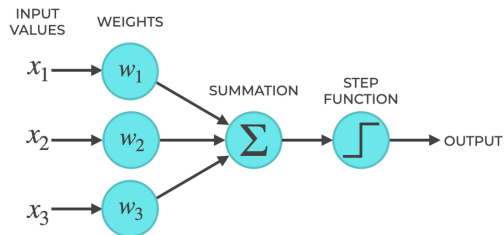
- 1 What is Deep Learning?
- 2 Brief History of Neural Networks
- 3 **Perceptron**
- 4 Multilayer Perceptrons

Biological Neuron



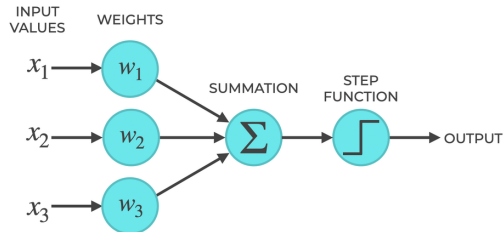
- **Dendrite:** Receives signals from other neurons
- **Soma:** Processes the information
- **Axon:** Transmits signals away from Soma
- **Axon terminal:** Send signals to other neurons

Perceptron



- Each input x_i is multiplied by its corresponding weight w_i
- The weighted inputs are summed together (along with the bias b).
- The sum is passed through a step function to produce the estimated output.

Mathematical Form of Perceptron



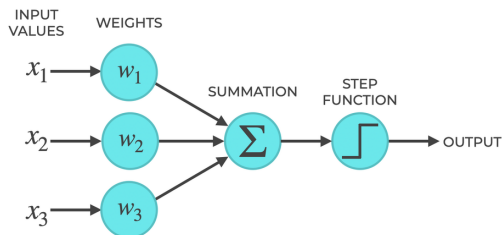
Mathematical Form of Perceptron:

$$\hat{y} = \phi \left(\sum_{i=1}^n w_i x_i + b \right)$$

- x_i are the **input**, w_i are the **weights**, b is the **bias**, \hat{y} is the **prediction**, ϕ is the **step function**:

$$\phi(z) = \begin{cases} 1, & \text{if } z \geq 0 \\ 0, & \text{if } z < 0 \end{cases}$$

Perceptron Example



Consider

- Inputs: $x_1, x_2 \in \{0, 1\}$ are **binary** values
- Weights: $w_1 = 1, w_2 = 1$
- Bias: $b = -1.5$

The perceptron output is computed as:

$$\hat{y} = \phi(w_1x_1 + w_2x_2 + b) = \phi(1 \cdot x_1 + 1 \cdot x_2 - 1.5)$$

- **Input:** $x_1 = 0, x_2 = 0$; **Output:** $z = 1 \cdot 0 + 1 \cdot 0 - 1.5 = -1.5 \Rightarrow \hat{y} = \phi(z) = 0$.
- **Input:** $x_1 = 0, x_2 = 1$; **Output:** $z = 1 \cdot 0 + 1 \cdot 1 - 1.5 = -0.5 \Rightarrow \hat{y} = \phi(z) = 0$.
- **Input:** $x_1 = 1, x_2 = 0$; **Output:** $z = 1 \cdot 1 + 1 \cdot 0 - 1.5 = -0.5 \Rightarrow \hat{y} = \phi(z) = 0$.
- **Input:** $x_1 = 1, x_2 = 1$; **Output:** $z = 1 \cdot 1 + 1 \cdot 1 - 1.5 = 0.5 \Rightarrow \hat{y} = \phi(z) = 1$.

Conclusion

The perceptron correctly implements the **AND** operator.

Matrix-Vector Representation of Perceptron

$$\hat{y} = \phi \left(\sum_{i=1}^n x_i w_i + b \right)$$

- Define vectors \mathbf{x} , $\mathbf{w} \in \mathbb{R}^n$:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_n \end{bmatrix}$$

- The perceptron can be defined in vector form:

$$\hat{y} = \phi(\mathbf{w}^\top \mathbf{x} + b),$$

where the inner or dot product of \mathbf{w} and \mathbf{x} is given by

$$\mathbf{w}^\top \mathbf{x} = \sum_{i=1}^n w_i x_i$$

Using Perceptrons to Implement Logical Operators

Logical operators such as AND (\wedge), OR (\vee), and NOT (\neg):

- \wedge : the result is true if both x_1 and x_2 are
- \vee : the result is true if either x_1 and x_2 is
- \neg : the result is true if the input is not

x_1	x_2	$x_1 \wedge x_2$	$x_1 \vee x_2$	$\neg x_1$
0	0	0	0	1
0	1	0	1	1
1	0	0	1	0
1	1	1	1	0

Table: Boolean table illustrating \wedge , \vee , and \neg operations with x_1 and x_2

Perceptron can be used to implement them:

$$\hat{y} = \phi(\mathbf{w}^\top \mathbf{x} + b) = \phi(w_1 x_1 + w_2 x_2 + b)$$

- \wedge : $\mathbf{w} = [1, 1]$ and $b = -1.5$
- \vee : $\mathbf{w} = [1, 1]$ and $b = 0.5$
- \neg : $w = -1$ and $b = 0.5$

Limitations of Perceptron

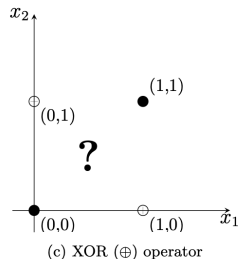
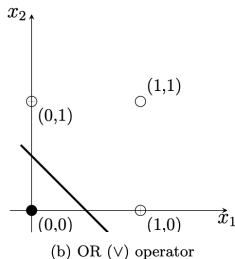
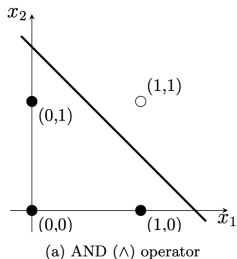
The perceptron cannot solve **nonlinear** problems such as the logical operator XOR (\oplus):

- \oplus : the result is true if x_1 and x_2 are different

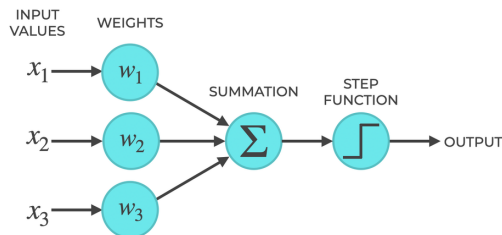
x_1	x_2	$x_1 \wedge x_2$	$x_1 \vee x_2$	$\neg x_1$	$x_1 \oplus x_2$
0	0	0	0	1	0
0	1	0	1	1	1
1	0	0	1	0	1
1	1	1	1	0	0

Table: Boolean table illustrating \wedge , \vee , and \oplus operations with x_1 and x_2

The perceptron can only solve **linearly** separable data:



Summary of Perceptron



The perceptron is defined as:

$$\hat{y} = \phi(\mathbf{w}^T \mathbf{x} + b),$$

- **Mathematical Model:** It computes the weighted sum of the inputs along with the bias term
- **Activation:** The perceptron (or neuron) is activated by the step (activation) function if the weighted sum exceeds a certain threshold.
- **Linear Separability:** The perceptron can classify linearly separable data, e.g., \wedge , \vee , and \neg .
- **Limitation:** It cannot solve **nonlinear** problems, e.g., \oplus .

- 1 What is Deep Learning?
- 2 Brief History of Neural Networks
- 3 Perceptron
- 4 **Multilayer Perceptrons**

Why Multilayer Perceptrons (MLP)?

Recap: A single perceptron can implement AND, OR, and NOT, but **not** XOR

However, XOR can be implemented by **multiple** perceptrons.

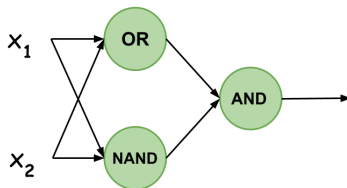
- **NAND** \uparrow : the result is false if both inputs are true

x_1	x_2	$x_1 \wedge x_2$	$x_1 \uparrow x_2$
0	0	0	1
0	1	0	1
1	0	0	1
1	1	1	0

- A single perceptron can implement NAND with $w = [-1, -1]$, and $b = 1.5$
- We can express XOR in terms of AND, OR, and NAND as follows

$$x_1 \oplus x_2 = (x_1 \vee x_2) \wedge (x_1 \uparrow x_2).$$

This forms a **2-layer network**:



MLP for XOR

The XOR function can be computed using MLP as follows:

- Define logical operators using single perceptrons:

$$h_i(\mathbf{x}) = \phi(\mathbf{w}_i^\top \mathbf{x} + b_i), \quad \forall i \in \{1, 2, 3\},$$

where the weight and biases are:

$$\text{OR} : \mathbf{w}_1 = [1, 1], \quad b_1 = -0.5$$

$$\text{NAND} : \mathbf{w}_2 = [-1, -1], \quad b_2 = 1.5$$

$$\text{AND} : \mathbf{w}_3 = [1, 1], \quad b_3 = -1.5$$

- Define the **activation vector** $\mathbf{a} \in \mathbb{R}^2$ as the intermediate results from the first layer:

$$\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \end{bmatrix}$$

where $a_1 = h_1(\mathbf{x})$ from the OR output and $a_2 = h_2(\mathbf{x})$ from the NAND output.

- Compute the estimated output \hat{y} using the activation vector \mathbf{a} from the first layer:

$$\hat{y} = h_3(\mathbf{a}) = \phi(\mathbf{w}_3^\top \mathbf{a} + b_3)$$

where h_3 represents the AND operation on the outputs of OR and NAND.

Structure of MLP

For each **hidden layer** in a multi-layer perceptron:

- The **input vector** $\mathbf{x} \in \mathbb{R}^d$ is obtained from the previous layer (or from the input data for the first layer).
- We define n perceptrons, each with independent **weights** $\mathbf{w}_i \in \mathbb{R}^d$ and a **bias** $b_i \in \mathbb{R}$ for $i \in [n] := \{1, 2, \dots, n\}$:

$$h_i(\mathbf{x}) = \phi(\mathbf{w}_i^\top \mathbf{x} + b_i)$$

where ϕ is the **activation function**.

- The computed outputs $h_i(\mathbf{x})$ are stacked into an **activation vector** $\mathbf{a} \in \mathbb{R}^n$, representing the output of this layer:

$$\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix}$$

where $a_i = h_i(\mathbf{x})$ for each i .

Matrix-Vector Representation of MLP

To rewrite the MLP in matrix-vector form for each layer ℓ :

- Define the **weight matrix** $\mathbf{W} \in \mathbb{R}^{n \times d}$ and the **bias vector** $\mathbf{b} \in \mathbb{R}^n$:

$$\mathbf{W} = \begin{bmatrix} \mathbf{w}_1^\top \\ \vdots \\ \mathbf{w}_n^\top \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} b_1 \\ \vdots \\ b_n \end{bmatrix}$$

where each $\mathbf{w}_i \in \mathbb{R}^d$ and b_i represents the weights and bias for the i -th perceptron.

- Define the **pre-activation vector** $\mathbf{z} \in \mathbb{R}^n$ as:

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b} = \begin{bmatrix} \mathbf{w}_1^\top \mathbf{x} + b_1 \\ \vdots \\ \mathbf{w}_n^\top \mathbf{x} + b_n \end{bmatrix}$$

This combines the weighted sum of the inputs from the previous layer.

- Apply the activation function ϕ **element-wise** to \mathbf{z} to obtain the activation vector $\mathbf{a} \in \mathbb{R}^n$:

$$\mathbf{a} = \phi(\mathbf{z}) = \begin{bmatrix} \phi(\mathbf{w}_1^\top \mathbf{x} + b_1) \\ \vdots \\ \phi(\mathbf{w}_n^\top \mathbf{x} + b_n) \end{bmatrix}$$

where ϕ is applied to each element of the pre-activation vector \mathbf{z} .

Summary of MLP

An MLP with L layers can be defined in a recurrent manner: for each layer $\ell \in [L]$,

$$\mathbf{z}^\ell = \mathbf{W}^\ell \mathbf{x}^{\ell-1} + \mathbf{b}^\ell,$$

$$\mathbf{x}^\ell = \phi(\mathbf{z}^\ell),$$

where $\mathbf{x}^\ell = \mathbf{a}^\ell$ serves as the input to the next layer, and the initial input is $\mathbf{x}^0 = \mathbf{x}$.

- ① The MLP is also called a **feed-forward network** because the data flows from the input layer to the output layer through hidden layers without any feedback loops.
- ② Each hidden layer consists of n_ℓ perceptrons (or neurons), where n_ℓ is referred to as the **width** of the network at layer ℓ .
- ③ The total number of layers L defines the **depth** of the network.
- ④ The final estimated output $\hat{\mathbf{y}} = \mathbf{x}^L$ can have **multiple dimensions**, depending on the task (e.g., classification or regression).
- ⑤ The **width** and **depth** are hyperparameters chosen by the network designer.
- ⑥ An MLP is capable of solving **nonlinear problems** that a single perceptron cannot handle.

Question

How do we effectively select the weights \mathbf{W}^ℓ and biases \mathbf{b}^ℓ ?