

# Optimization in Neural Networks

**Tianxiang (Adam) Gao**

September 23, 2024

# Outline

- 1 Calculus Review: Extreme Values
- 2 Convergence Issues
- 3 Advanced Optimization Algorithm

## Recap: Neural Network Training

We use a **training process** iteratively update the parameters in MLPs:

- MLPs are **parameterized** function  $f_{\theta}$ , where  $\theta = \{\mathbf{W}^{\ell}, \mathbf{b}^{\ell}\}$
- **Universal Approximation Theorem (UAT)**: MLPs can approximate “any” function  $f^*$  arbitrarily accurate, provided sufficient parameters (and training samples).
- Given a **training set**  $\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^{\ell}$  and a **loss** function  $\ell$ , the training problem is:

$$\min_{\theta} \quad \mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f_{\theta}(\mathbf{x}_i), \mathbf{y}_i)$$

- This optimization problem can be solved using **gradient descent**, which gradually reduces the cost  $\mathcal{L}$  along the *steepest descent direction*:

$$\theta^+ = \theta - \eta \nabla \mathcal{L}(\theta)$$

where  $\eta > 0$  is the **learning rate**.

- The gradients in MLPs can be computed using the **chain rule** backward from the total cost.

## Recap: Neural Network Training

- Using the **computational graph**, the gradients can be computed through **backpropagation**:

- Forward Propagation (biases omitted): Start with  $\mathbf{x}^0 = \mathbf{x}$

$$\mathbf{z}^\ell = \mathbf{W}^\ell \mathbf{x}^{\ell-1}, \quad \forall \ell \in \{0, 1, 2, \dots, L\}$$

$$\mathbf{x}^\ell = \phi(\mathbf{z}^\ell),$$

- Backward Propagation (biases omitted): Start with  $d\mathbf{z}^L = (\mathbf{x}^L - \mathbf{y}) \odot \phi'(\mathbf{z}^L)$

$$d\mathbf{z}^\ell = \left[ (\mathbf{W}^{\ell+1})^\top d\mathbf{z}^{\ell+1} \right] \odot \phi'(\mathbf{z}^\ell), \quad \forall \ell \in \{1, 2, \dots, L-1\}$$

$$d\mathbf{W}^\ell = d\mathbf{z}^\ell \mathbf{x}^{(\ell-1)\top}$$

- Random initialization** is preferred over zero initialization to avoid the issue of *symmetric patterns*.

### Questions


- How do I choose the right learning rate?
- How do I determine the appropriate width and depth of the network?
- Does gradient descent (GD) always converge?
- How can I speed up the training process?

# Outline

- 1 Calculus Review: Extreme Values
- 2 Convergence Issues
- 3 Advanced Optimization Algorithms

# Calculus Review: Extreme Values

Let  $f(x)$  be a real-valued function, where  $x \in \mathbb{R}$ .




Local Min.  $x_1, x_3, x_5, x_7, x_9$ ;

Local Max.  $x_2, x_4, x_6, x_8$ ;

- The function  $f$  has an **local minimum** at point  $x = a$  if  $f(a) \leq f(x)$  when  $x$  is near  $a$ .
- The function  $f$  has an **local maximum** at point  $x = a$  if  $f(a) \geq f(x)$  when  $x$  is near  $a$ .
- The point  $a$  is a **global minimum** or **global maximum** if the above property holds for all  $x$ .
- **Fermat's Theorem:** If  $f$  has a local min or max at  $x = a$ , then  $f'(a) = 0$
- A point  $x = a$  is called **stationary** if  $f'(a) = 0$ .
- **Gradient descent** stops at *stationary points*:

$$\theta^+ = \theta - \eta \nabla_{\theta} \mathcal{L}(\theta).$$

# Calculus Review: Curvature



**Concavity:** the second derivative  $f''$  describes whether  $f$  is **concave up** or **concave down**

- If  $f'' > 0$ , then  $f$  is **concave up** at  $x$ .
- If  $f'' < 0$ , then  $f$  is **concave down** at  $x$ .

**The Second Derivative Test:**

- If  $f'(a) = 0$  and  $f''(a) \geq 0$ , then  $a$  is a **local minimum**
- If  $f'(a) = 0$  and  $f''(a) \leq 0$ , then  $a$  is a **local maximum**.

## Conclusion

The goal of training in deep learning is to find a good **local minimum** that generalizes well.

# Hessian Matrix

Let  $f(x)$  be a **multivariate** real-valued function, where  $x \in \mathbb{R}^n$ .

- A point  $x = a$  is called **stationary point** if  $\nabla f(a) = \mathbf{0}$ , i.e.,

$$\nabla f(a) = \left[ \frac{\partial f(a)}{\partial x_1} \quad \cdots \quad \frac{\partial f(a)}{\partial x_n} \right]^\top = \mathbf{0}$$




- The **Hessian** matrix  $H(w) \in \mathbb{R}^{n \times n}$  of  $f$  is the symmetric matrix of second-order partial derivatives:

$$\nabla^2 f(x) = H(x) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

- For the **second-order mixed partial derivative**  $\frac{\partial^2 f}{\partial x \partial y}$  is the rate of change of  $\frac{\partial f}{\partial x}$  w.r.t.  $y$  changes, holding  $x$  constant.



# Significance of Hessian



## Interpretation of the Hessian Matrix:

- The Hessian describes the **local curvature** of the function.
- **Positive** definite Hessian  $H$  implies a local minimum, *i.e.*, concave up in any direction.
- **Negative** definite Hessian implies a local maximum, *i.e.*, concave down in any direction.
- **Indefinite** Hessian implies a **saddle point**, *i.e.*, concave up in some directions and concave down in others.

# Discussion Questions

Compute the gradients and Hessian of the following functions:

- $f(x) = \frac{1}{2}(xw - y)^2$
- $f(w) = \frac{1}{2}\|Xw - y\|^2$ , where

$$w = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}, \quad X = \begin{bmatrix} 3 & 1 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

**Instructions:** Discuss these questions in small groups of 2-3 students.

# Solutions to the Discussion Questions

Compute the gradients and Hessian of the following functions:

- $f(x) = \frac{1}{2}(xw - y)^2$ ,  $f'(w) = x \cdot (xw - y)$ , and  $f''(w) = x^2$ .
- $f(\mathbf{w}) = \frac{1}{2}\|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$ , where

$$\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}, \quad \mathbf{X} = \begin{bmatrix} 3 & \\ & 1 \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

We have

$$\nabla f(\mathbf{w}) = \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) = \begin{bmatrix} 3(3w_1 - 1) \\ w_2 \end{bmatrix} \quad \text{and} \quad \mathbf{H}(\mathbf{w}) = \begin{bmatrix} 9 & \\ & 1 \end{bmatrix},$$



- Here 9 is the **largest eigenvalue** of  $\mathbf{H}$ , 1 is the **smallest eigenvalue** of  $\mathbf{H}$ , and their ratio is called **conditional number**  $\kappa = 9$ .

# Outline

- 1 Calculus Review: Extreme Values
- 2 **Convergence Issues**
- 3 Advanced Optimization Algorithm

## Learning Rate

## Learning Rate



# One-Dimensional Linear Regression

Consider a simple one-dimensional linear regression problem:

$$\min_w \mathcal{L}(w) = \ell(f_\theta(x), y) = \frac{1}{2}(wx - y)^2,$$

where  $w, x, y \in \mathbb{R}$ .

- The function  $f_\theta(x) = wx$  is a perceptron with linear activation, without a bias term.
- With gradient  $\nabla \mathcal{L}(w) = x(wx - y)$ , the gradient descent update is:

$$w^+ = w - \eta \cdot x(wx - y),$$

where  $\eta > 0$  is the learning rate.

- To find the **stationary point**:

$$\nabla \mathcal{L}(w) = 0 \implies x(wx - y) = 0 \implies w^* = \frac{y}{x}$$

- Second derivative test:

$$\nabla^2 \mathcal{L}(w^*) = x^2 > 0,$$

*i.e.*,  $w^*$  is a local minimum (and also a global minimum since  $\mathcal{L}$  is concave up everywhere).

# Recursive Formula for Gradient Descent on LSR

- The update rule for Gradient Descent applied to linear regression is:

$$w^{k+1} = w^k - \eta \cdot x(w^k x - y) = (1 - \eta x^2)w^k + \eta xy := aw^k + b,$$

where  $a := 1 - \eta x^2$  and  $b := \eta xy$ .

- Using this recurrence relation,  $w^{k+1}$  can be expanded as:

$$\begin{aligned} w^{k+1} &= aw^k + b \\ &= a(aw^{k-1} + b) + b \\ &= a^2w^{k-1} + ab + b \\ &= a^3w^{k-2} + a^2b + ab + b \\ &= a^{k+1}w^0 + b(a^k + a^{k-1} + \cdots + a + 1) \\ &= a^{k+1}w^0 + b \frac{1 - a^{k+1}}{1 - a} \\ &= a^{k+1}(w^0 - w^*) + w^*, \end{aligned}$$


where we use the geometric series  $\sum_{i=0}^k a^i = \frac{1 - a^{k+1}}{1 - a}$  and  $w^* = \frac{y}{x}$ .

# Impact of Learning Rate on Convergence


The recurrence relation:

$$w^{k+1} = a^{k+1}(w^0 - w^*) + w^*,$$


where  $a = 1 - \eta x^2$ , leads to the following behaviors as  $k \rightarrow \infty$ :




Slow



Just Right



Oscillation



Divergence

- **Convergence:** If  $\eta < 2/x^2$ , then  $|a| < 1$ , so  $a^k \rightarrow 0$ , and  $w^k$  converges to the minimum  $w^*$ .
- **Oscillation:** If  $\eta = 2/x^2$ , then  $a = -1$ , and  $w^k$  oscillates around  $w^*$  with  $w^{k+1} = (-1)^{k+1}(w^0 - w^*) + w^*$ .
- **Divergence:** If  $\eta > 2/x^2$ , then  $|a| > 1$ , leading to  $a^k \rightarrow \infty$ , causing  $w^k$  to diverge.



# Residual Dynamics in Gradient Descent

The update rule for Gradient Descent on LSR is:

$$w^{k+1} = w^k - \eta \cdot x(w^k x - y).$$

From this, we can derive a recurrence relation for the residual or error  $\varepsilon^{k+1}$ :

$$\begin{aligned}\varepsilon^{k+1} &= w^{k+1}x - y \\ &= \left[ w^k - \eta \cdot x(w^k x - y) \right] x - y \\ &= (1 - \eta x^2) \cdot \varepsilon^k \\ &= a \cdot \varepsilon^k,\end{aligned}$$


where  $a := 1 - \eta x^2$  and  $\varepsilon^k = w^k x - y$  is the error at step  $k$ .

Repeating this relation, we obtain:

$$\varepsilon^{k+1} = a^{k+1} \varepsilon^0,$$

where  $\varepsilon^0 = w^0 x - y$  is the initial error.

## Loss Landscape



# Curse of Dimensionality in Optimization

- As the dimensionality of variables and the size of data increase, optimization becomes more challenging. For example, consider the following loss function:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \frac{1}{2} (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 = \frac{1}{2n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

- The recurrence relation for  $\mathbf{w}^{k+1}$  becomes:



$$\mathbf{w}^{k+1} = \left( \mathbf{I} - \frac{\eta}{n} \mathbf{X} \mathbf{X}^\top \right)^{k+1} (\mathbf{w}^0 - \mathbf{w}^*) + \mathbf{w}^* = \mathbf{A}^{k+1} (\mathbf{w}^0 - \mathbf{w}^*) + \mathbf{w}^*,$$

where  $\mathbf{A} := \mathbf{I} - \frac{\eta}{n} \mathbf{X} \mathbf{X}^\top$ .


- The dynamics are governed by the matrix  $\mathbf{A}$ , rather than a scalar. In deep learning, this system becomes even more complex as  $\mathbf{A}$  can change during training, making it difficult to maintain a clear recurrence structure.

## 3D Loss Landscape Visualization

Consider a case where  $\mathbf{w} = (w_1, w_2)$ . Below is the 3D contour of  $\mathcal{L}(\mathbf{w})$ :



Well-Conditioned




Ill-Conditioned

The loss landscape is not always smooth and easy to optimize:


$$\mathbf{X} = \begin{bmatrix} 1 & 0.1 \\ 0 & 1 \end{bmatrix}, \quad \mathbf{X} = \begin{bmatrix} 3 & 0.1 \\ 0 & 1 \end{bmatrix}$$

# Challenges in Gradient Descent: Zig-Zag Patterns


- In ill-conditioned systems, gradient descent can only progress with a small learning rate. The following examples illustrate different behaviors:



**Fast Convergence ( $\eta = 1.0$ )**



**Divergence ( $\eta = 0.23$ )**



**Zig-Zag Pattern ( $\eta = 0.22$ )**

## Key Observations

- Ill-conditioned systems cannot tolerate large learning rates.
- Even with a small learning rate, gradient descent may exhibit a zig-zag pattern.

## III-Conditioned Systems

Consider the recurrence relation for ill-conditioned systems:

$$\begin{aligned} \mathbf{w}^{k+1} &= \left( \mathbf{I} - \frac{\eta}{n} \mathbf{X} \mathbf{X}^\top \right)^{k+1} (\mathbf{w}^0 - \mathbf{w}^*) + \mathbf{w}^* \\ &= \begin{bmatrix} 1 - \frac{9\eta}{2} & \\ & 1 - \frac{\eta}{2} \end{bmatrix}^{k+1} (\mathbf{w}^0 - \mathbf{w}^*) + \mathbf{w}^*. \end{aligned}$$

where we use  $n = 2$  and

$$\mathbf{X} = \begin{bmatrix} 3 & \\ & 1 \end{bmatrix}, \quad \text{and} \quad \mathbf{X} \mathbf{X}^\top = \begin{bmatrix} 9 & \\ & 1 \end{bmatrix}$$

- From the first exponential, convergence requires  $\eta < \frac{4}{9}$ .
- From the second exponential, convergence requires  $\eta < 4$ .

### Key Observations: Condition Number and Learning Rate

- To ensure convergence, we must choose  $\eta < \frac{4}{9}$ .
- One direction converges much slower than the other, leading to the zig-zag behavior.
- This occurs because the condition number  $\kappa$  of the Hessian  $\mathbf{H}(\mathbf{w})$  is large, i.e.,  $\kappa = 9$ .

## Gradients Vanishing and Exploding

# Gradients Vanishing and Exploding

# Information Propagation in Deep Neural Networks

- **Forward Propagation (biases omitted):** Starting with  $\mathbf{x}^0 = \mathbf{x}$ ,

$$\mathbf{z}^\ell = \mathbf{W}^\ell \mathbf{x}^{\ell-1}, \quad \forall \ell \in \{0, 1, 2, \dots, L\},$$

$$\mathbf{x}^\ell = \phi(\mathbf{z}^\ell),$$

where  $\phi(\mathbf{z})$  is the activation function.

- Assuming a linear activation function  $\phi(\mathbf{z}) = \mathbf{z}$  for simplicity:

$$\mathbf{x}^\ell = \mathbf{W}^\ell \mathbf{x}^{\ell-1} = \begin{bmatrix} a & \\ & a \end{bmatrix}^\ell \mathbf{x}^0 = a^\ell \mathbf{x}^0.$$

As  $\ell$  increases:

- If  $a > 1$ , then  $\mathbf{x}^\ell$  grows exponentially (explodes).
- If  $a < 1$ , then  $\mathbf{x}^\ell$  diminishes exponentially (vanishes).



# Backward Propagation and Gradient Behavior

- **Backward Propagation (biases omitted):** Start with  $d\mathbf{z}^L = (\mathbf{x}^L - \mathbf{y}) \odot \phi'(\mathbf{z}^L)$ :

$$\begin{aligned} d\mathbf{z}^\ell &= \left[ (\mathbf{W}^{\ell+1})^\top d\mathbf{z}^{\ell+1} \right] \odot \phi'(\mathbf{z}^\ell), \quad \forall \ell \in \{1, 2, \dots, L-1\}, \\ d\mathbf{W}^\ell &= d\mathbf{z}^\ell \mathbf{x}^{(\ell-1)\top}. \end{aligned}$$

- With linear activation,  $\phi'(x) = 1$ :

$$d\mathbf{z}^\ell = (\mathbf{W}^{\ell+1})^\top d\mathbf{z}^{\ell+1} = \begin{bmatrix} a & \\ & a \end{bmatrix} d\mathbf{z}^{\ell+1} = a^{L-\ell} d\mathbf{z}^L.$$

- As  $\ell$  becomes far from  $L$ :
  - If  $a > 1$ , then  $d\mathbf{z}^\ell$  grows rapidly (exploding gradients).
  - If  $a < 1$ , then  $d\mathbf{z}^\ell$  diminishes rapidly (vanishing gradients).

# Summary

## Learning Rate:

- Small learning rates slow down the training.
- Large learning rates can cause oscillations or divergence.

## Loss Landscape:

- The loss landscape is often ill-conditioned in DNNs, with local minima, maxima, and saddle points.
- Ill-conditioned local structure prevents using a large learning rate.

## Gradient Vanishing and Exploding:

- Information propagation in DNNs can be unstable.
- Lower layers tend to have small gradient values due to vanishing gradients.
- Differing gradient scales can lead to ill-conditioned local structures.



## Outline

- 1 Calculus Review: Extreme Values
- 2 Convergence Issues
- 3 **Advanced Optimization Algorithm**

## Gradient Descent with Momentum

# The Trajectory of Gradient Descent

Let us take a close look at the trajectory of gradient descent (GD):



## Average Search Direction

- The **general** iterative training process is defined as:


$$\mathbf{w}^+ = \mathbf{w} - \eta \cdot \mathbf{v},$$

where  $\mathbf{v}$  is the search direction, and  $\eta$  is the learning rate. We take  $\mathbf{v} = \nabla \mathcal{L}(\mathbf{w})$  for GD.

- Given a trajectory of GD up to the  $k$ -th iteration, the sequence of gradient directions is:

$$\{\mathbf{g}^0, \mathbf{g}^1, \dots, \mathbf{g}^{k-1}\}, \quad \text{where } \mathbf{g}^i = \nabla \mathcal{L}(\mathbf{w}^i) \implies \mathbf{v}^k = \frac{1}{k} \sum_{i=0}^{k-1} \mathbf{g}^i.$$

- Smooth out noisy gradients and maintain a more stable descent trend over iterations



## GD with Averaged Gradient Direction

By applying the idea of averaging the negative gradient direction, we have:

$$\begin{aligned} \mathbf{v}^{k+1} &= \frac{1}{k+1} \sum_{i=0}^k \mathbf{g}^i, \\ \mathbf{w}^{k+1} &= \mathbf{w}^k - \eta \cdot \mathbf{v}^{k+1}. \end{aligned}$$

- The **cumulative average** can be rewritten in a **running update** form:

$$\begin{aligned} \mathbf{v}^{k+1} &= \frac{1}{k+1} \left( \sum_{i=0}^{k-1} \mathbf{g}^i + \mathbf{g}^k \right) \\ &= \frac{k}{k+1} \cdot \frac{1}{k} \sum_{i=0}^{k-1} \mathbf{g}^i + \frac{1}{k+1} \mathbf{g}^k \\ &= \frac{k}{k+1} \mathbf{v}^k + \left( 1 - \frac{k}{k+1} \right) \mathbf{g}^k. \end{aligned}$$

- Hence, gradient descent with an averaged gradient direction is given by:

$$\begin{aligned} \mathbf{v}^{k+1} &= \beta_{k+1} \mathbf{v}^k + (1 - \beta_{k+1}) \mathbf{g}^k, \\ \mathbf{w}^{k+1} &= \mathbf{w}^k - \eta \cdot \mathbf{v}^{k+1}, \end{aligned}$$

where  $\beta_{k+1} = \frac{k}{k+1}$ .

# Gradient Descent with Momentum

- Fixing  $\beta_{k+1} = \beta$  for  $\beta \in (0, 1)$ , i.e.,  $\beta = 0.9$ , the update rule becomes:

$$\mathbf{v}^{k+1} = \beta \mathbf{v}^k + (1 - \beta) \mathbf{g}^k,$$

$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \cdot \mathbf{v}^{k+1},$$

- Here,  $\beta$  controls how much influence past gradients ( $\mathbf{v}^k$ ) have on the current update versus the most recent gradient ( $\mathbf{g}^k$ ).
- This method is also referred to as **Gradient Descent (GD) with Momentum** or **accelerated GD**:

$$\begin{aligned} \mathbf{w}^{k+1} &= \mathbf{w}^k - \eta \cdot \mathbf{v}^{k+1} \\ &= \mathbf{w}^k - \eta \cdot [\beta \mathbf{v}^k + (1 - \beta) \mathbf{g}^k] \\ &= \mathbf{w}^k - \eta(1 - \beta) \cdot \mathbf{g}^k + \beta \cdot \underbrace{(\mathbf{w}^k - \mathbf{w}^{k-1})}_{\text{Momentum}}, \quad \text{as } \mathbf{w}^k = \mathbf{w}^{k-1} - \eta \cdot \mathbf{v}^k. \end{aligned}$$



The current update is influenced **both** by the latest gradient and the past movement (momentum).



# Impact of Momentum in Gradient Descent

Gradient Descent with momentum:

$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta(1 - \beta) \cdot \mathbf{g}^k + \beta \cdot (\mathbf{w}^k - \mathbf{w}^{k-1}).$$






- Gradient Descent (GD) converges in 84 steps with  $\eta = 0.22$ , while GD with momentum converges in 36 steps with  $\eta = 0.63$  and  $\beta = 0.6$ .
- For further reading, see this [illustration on the impact of momentum](#).

# Damping in Gradient Descent with Momentum

Gradient Descent with momentum:

$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \cdot (1 - \beta) \cdot \mathbf{g}^k + \beta \cdot (\mathbf{w}^k - \mathbf{w}^{k-1}).$$



## Key Observation

- A large momentum factor  $\beta$  can cause the loss to oscillate and not consistently decrease.
- This oscillation often occurs around the stationary point.

# Summary

Gradient Descent with momentum:



$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta(1 - \beta) \cdot \mathbf{g}^k + \beta \cdot (\mathbf{w}^k - \mathbf{w}^{k-1}).$$

- The current update is influenced by **both** the most recent gradient and the past movement.
- The search direction in GD with momentum is a **running average** of past gradients.
- Momentum allows for **larger learning rates** and **faster** convergence.
- Too large a momentum factor  $\beta$  can cause **damping** in the loss and oscillation around the stationary point.

# Adaptive Gradient Descent

# Divergent Gradient Scaling

During the GD, the **magnitudes** of the gradient coordinates can vary significantly. One approach is to **scale** the magnitudes so that each gradient coordinate has an order of  $\mathcal{O}(1)$  magnitude.



## RMSProp

- By applying the idea of a *running average* on the **gradient magnitudes**, the scaling factors are:

$$\begin{aligned}s^+ &= \beta s + (1 - \beta)g^{\odot 2}, \\ w^+ &= w - \eta \cdot g \oslash \sqrt{s^+ + \varepsilon},\end{aligned}$$

where  $\odot$  denotes element-wise multiplication,  $\sqrt{\cdot}$  represents the element-wise square root,  $\oslash$  denotes element-wise division, and  $\varepsilon$  is a small value (e.g.,  $\varepsilon = 10^{-8}$ ) preventing dividing by zero.

- This method is called **root mean squared propagation (RMSP)**.
- RMSProp is effectively an **adaptive learning rate** algorithm:

$$w_i^+ = w_i - \eta_i \frac{\partial \mathcal{L}(w)}{\partial w_i},$$

where  $\eta_i = \frac{\eta}{\sqrt{s_i^+}}$  is the **adaptive learning rate**.



- Each gradient coordinate has a unique, adaptive learning rate.

# Performance of RMSProp

RMSProp:

$$\mathbf{s}^+ = \beta \mathbf{s} + (1 - \beta) \mathbf{g}^{\odot 2},$$

$$\mathbf{w}^+ = \mathbf{w} - \eta \cdot \mathbf{g} \oslash \sqrt{\mathbf{s}^+}.$$






- GD converges in 84 steps with  $\eta = 0.22$ .
- RMSProp converges in 43 steps with  $\eta = 0.07$  and in 10 steps with  $\eta = 0.22$ .
- **Note:** RMSProp may **not** perform well with large learning rates.

## Adam

The **Adaptive Moment Estimation (Adam)** algorithm combines the advantages of GD with momentum and RMSProp:

$$\begin{aligned} \mathbf{v}^+ &= \beta_1 \mathbf{v} + (1 - \beta_1) \mathbf{g}, \\ \mathbf{s}^+ &= \beta_2 \mathbf{s} + (1 - \beta_2) \mathbf{g}^{\odot 2}, \\ \mathbf{w}^+ &= \mathbf{w} - \eta \cdot \mathbf{v}^+ \oslash \sqrt{\mathbf{s}^+}, \end{aligned}$$

where typical values in training DNNs are  $\beta_1 = 0.9$  and  $\beta_2 = 0.99$ .



- GD converges in 84 steps with  $\eta = 0.22$ .
- GD with momentum converges in 35 steps with  $\eta = 0.63$  and  $\beta = 0.6$ .
- RMSProp converges in 43 steps with  $\eta = 0.07$ .
- Adam converges in 32 steps with  $\eta = 0.74$ .



# Summary

- Adaptive gradient descent (AdaGrad) scales each gradient coordinate to have  $\mathcal{O}(1)$  magnitudes.
- Adaptive methods provide an **adaptive learning rate** for each gradient coordinate.
- Typically, adaptive methods do not use large learning rates.
- Adam is a combination of momentum-based and adaptive scaling techniques, balancing fast convergence with gradient smoothing.

# Stochastic Gradient Descent

# Stochastic Gradient Descent (SGD) Overview

**Recap:** Training deep neural networks as an optimization problem over parameters  $\theta$ :

$$\min_{\theta} \mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f_{\theta}(\mathbf{x}_i), y_i)$$

- The gradient descent (GD) update rule is:

$$\theta^+ = \theta - \eta \nabla_{\theta} \mathcal{L}(\theta) = \theta - \eta \cdot \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} \ell_i(\theta),$$

where  $\ell_i(\theta) := \ell(f_{\theta}(\mathbf{x}_i), y_i)$  is the loss for sample  $i$ .

- In practice, the number of training samples  $n$  can be extremely large (millions or even billions). Computing the gradient over all samples becomes computationally expensive.
- **Stochastic Gradient Descent (SGD):** Instead of computing the gradient over the full dataset, we randomly select a smaller batch  $\mathcal{B}$  of samples (called a **mini-batch**):

$$\theta^+ = \theta - \eta \cdot \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \nabla_{\theta} \ell_i(\theta)$$

- The size of the mini-batch  $|\mathcal{B}|$  can vary. If  $|\mathcal{B}| = 1$ , it is called **SGD**. Otherwise, it is called **mini-batch SGD**.

# Mini-batch SGD and Epochs

- In mini-batch SGD, the entire dataset is typically divided into several mini-batches of a fixed size  $b$ .
- The mini-batches are often selected by random shuffling (or **permutation**), and the model is updated iteratively for each mini-batch.
- After processing all mini-batches once, we complete an **epoch**, and the process can be repeated for multiple epochs until convergence.
- **Efficiency:** Mini-batch SGD can be computationally efficient because each update is based on a subset of data, reducing the cost per iteration.
- **Advanced Techniques:** Mini-batch SGD can be combined with other optimization techniques, such as momentum, RMSProp, and Adam.

# SGD vs. Full Batch Gradient Descent

- **Stochastic Behavior:** Unlike full-batch gradient descent, the loss function in SGD does not always decrease at every step due to the randomness of mini-batches. This can cause oscillations.
- **Convergence Speed:** Although SGD may take more iterations to converge in theory, it often converges faster in terms of wall-clock time due to its lower per-iteration computational cost.
- **Trade-off:** Full-batch GD ensures a consistent reduction in loss at each step, but the cost per iteration is high, especially for large datasets. SGD trades off some accuracy for faster convergence.